

In this lab you implement/simulate the scheduling and optimization of I/O operations. Applications submit their IO requests to the IO subsystem (potentially via the filesystem and pagecache), where they are maintained in an IO-queue until the disk device is ready for servicing another request. The IO-scheduler then selects a request from the IO-queue and submits it to the disk device. This selection is commonly known as the `strategy()` routine in operating systems. On completion, another request can be taken from the IO-queue and submitted to the disk. The scheduling policies will allow for some optimization as to reduce disk head movement or overall wait time in the system. The schedulers to be implemented are FIFO (i), SSTF (j), LOOK (s), CLOOK (c), and FLOOK (f) (the letters in bracket define which parameter must be given in the `-s` program flag).

You are to implement these different IO-schedulers in C or C++ and submit the **source** code, which we will compile and run. Your submission must contain a Makefile so we can run on **linserv*.cims.nyu.edu** (and please at least test there as well).

Invocation is as follows:

```
./iosched -s<schedalgo> [options] <inputfile> // options are -v -q -f (see below)
```

The input file is structured as follows: Lines starting with '#' are comment lines and should be ignored.

Any other line describes an IO operation where the 1st integer is the time step at which the IO operation is issued and the 2nd integer is the track that is accesses. Since IO operation latencies are largely dictated by seek delay (i.e. moving the head to the correct track), we ignore rotational and transfer delays for simplicity. The inputs are well formed.

```
#io generator
#numio=32 maxtracks=512 lambda=10.000000
1 430
129 400
:
```

We assume that moving the head by one track will cost one time unit. As a result, your simulation can/should be done using integers. The disk can only consume/process one IO request at a time. Everything else must be maintained in an IO queue and managed according to the scheduling policy. The initial direction of the LOOK algorithms is from 0-tracks to higher tracks. The head is initially positioned at track=0 at time=0. Note that you do not have to know the maxtrack (think SCAN vs. LOOK).

Each simulation should print information on individual IO requests followed by a SUM line that has computed some statistics of the overall run. (see reference outputs).

For each IO request create an info line (5 requests shown).

```
0:      1      1    431
1:     87    467    533
2:    280    431    467
3:    321    533    762
4:    505    762    791
```

Created by

```
printf("%5d: %5d %5d %5d\n",i, req->arrival_time, r->start_time, r->end_time);
- IO-op#,
- its arrival to the system (same as inputfile)
- its disk service start time
- its disk service end time
```

Please remember “ %5d” is not “%6d” !!!

and for the complete run a SUM line:

total_time: total simulated time, i.e. until the last I/O request has completed.
tot_movement: total number of tracks the head had to be moved
avg_turnaround: average turnaround time per operation from time of submission to time of completion
avg_waittime: average wait time per operation (time from submission to issue of IO request to start disk operation)
max_waittime: maximum wait time for any IO operation.

```
Created by: printf("SUM: %d %d %.2lf %.2lf %d\n",  
                  total_time, tot_movement, avg_turnaround, avg_waittime, max_waittime);
```

Various sample inputs and outputs are provided with the assignments on NYU classes.
Please look at the sum results and identify what different characteristics the schedulers exhibit.

You can make the following assumptions:

- at most 10000 IO operations will be tested, so its OK (recommended) to first read all requests from file before processing.
- all io-requests are provided in increasing time order (no sort needed)
- you never have two IO requests arrive at the same time (so input is monotonically increasing)

I strongly suggest, you do not use discrete event simulation this time. You can write a loop that increments simulation time by one and checks whether any action is to be taken. In that case you have to check in the following order.

- 1) If a new I/O arrived to the system at this current time → add request to IO-queue ?
- 2) If an IO is active and completed at this time → Compute relevant info and store in IO request for final summary
- 3) If an IO is active but did not yet complete → Move the head by one sector/track/unit in the direction it is going (to simulate seek)
- 4) If no IO request active now (after (2)) but IO requests are pending → Fetch the next request and start the new IO.
- 5) If no IO request is active now and no IO requests pending → exit simulation
- 6) Increment time by 1

When switching queues in FLOOK you always continue in the direction you were going from the current position, until the queue is empty. Then you switch direction until empty and then switch the queues continuing into that direction and so forth. While other variants are possible, I simply chose this one this time though other variants make also perfect sense.

Additional Information:

As usual, I provide some more detailed tracing information to help you overcome problems. Note your code only needs to provide the result line per IO request and the ‘SUM line’.

The reference program under ~frankeh/Public/iosched on the cims machine implements three options -v, -q, -f to debug deeper into IO tracing and IO queues.

The -v execution trace contains 3 different operations (add a request to the IO-queue, issue an operation to the disk and finish a disk operation). Following is an example of tracking IO-op 18 through the times 1139..1295 from submission to completion.

```
1139: 18 add 211          // 18 is the IO-op # (starting with 0) and 211 is the track# requested  
1127: 18 issue 211 279    // 18 is the IO-op #, 211 is the track# requested, 279 is the current track#  
1195: 18 finish 68        // 18 is the IO-op #, 68 is total length/time of the io from request to completion
```

-q shows the details of the IO queue and direction of movement (1==up , -1==down) and
-f shows additional queue information during the FLOOK.

Here Queue entries are tuples during add [ior# : #io-track] or triplets during get [ior# : io-track# : distance], where distance is negative if it goes into the opposite direction (where applicable).

Please use these debug flags and the reference program to get more insights on debugging the ins and outs (no punt intended) of this assignment and answering certain “why” questions.